

Predictive Modeling

XGBoost & CatBoost

DS Development Presentations
Peter Nicholas S. Onglao | MNL

- Guide to XGBoost Hyperparameters
- Options for Hyperparameter Tuning
- Introduction to CatBoost

1

XGB Hyperparameters

A guide to relevant parameters



List of Parameters

Learning Task Parameters		
Parameter	Default	Description
objective	"reg:squarederror"	The objective function ex. reg:logistic, binary:logistic, multi:softprob
base_score	0.5	Initial prediction score of all instances
eval_metric	Depends on obj.	Evaluation metric for validation data
seed	0	Random seed number (use set.seed() instead in R)



List of Parameters

Parameters for Tree Booster		
Parameter	Default	Description
eta	0.3	Step size shrinkage to prevent overfitting
gamma	0	Minimum loss reduction to make a further partition
lambda	1	L2 regularization term on weights
alpha	0	L1 regularization term on weights

$$Gain = \left[\frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{(G_L + G_R)^2}{H_L + H_R + \lambda} \right] - \gamma$$

$$\frac{G^2}{H + \lambda} = \frac{(\sum g)^2}{\sum h + \lambda} = \frac{(\sum (residuals))^2}{\sum \hat{y}_i \times (1 - \hat{y}_i) + \lambda}$$



List of Parameters

Parameters for Tree Booster		
Parameter	Default	Description
max_depth	6	Maximum tree depth
min_child_weight	1	Minimum 'score' needed in a child; further partitioning stops if a tree partition results in a node less than this value
max_delta_step	0	Usually not needed, but helps in logistic regression when class is extremely imbalanced. Set between 1-10

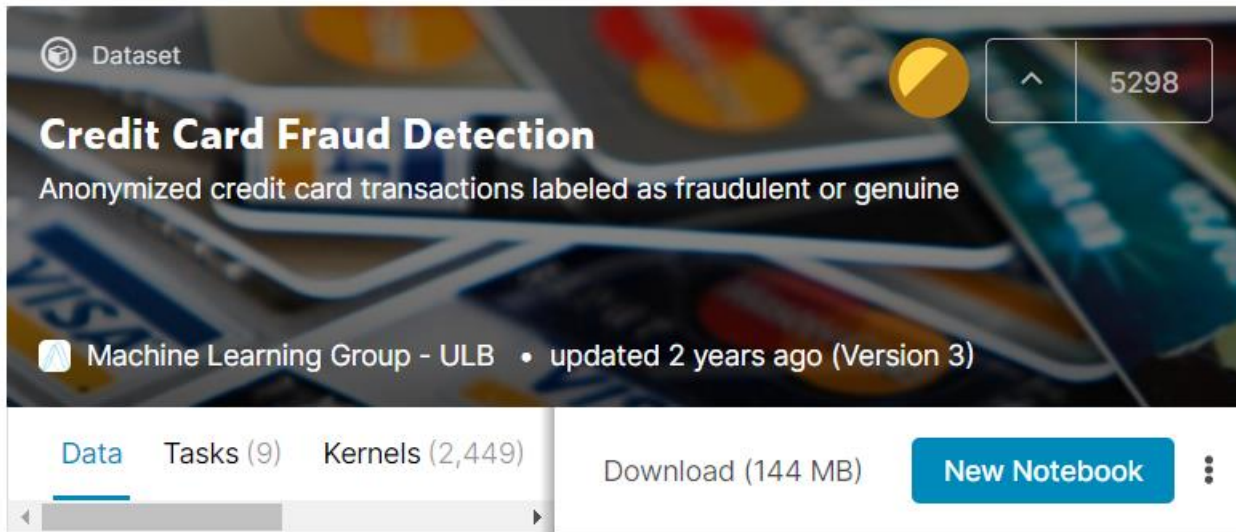
$$\frac{G^2}{H + \lambda} = \frac{(\sum g)^2}{\sum h + \lambda} = \frac{(\sum residuals)^2}{\sum \hat{y}_i \times (1 - \hat{y}_i) + \lambda}$$






List of Parameters

Parameters for Tree Booster		
Parameter	Default	Description
subsample	1 range: (0,1]	Subsamples training data at every iteration
colsample_bytree	1 range: (0,1]	Subsamples columns for each tree
colsample_bylevel		Subsamples columns for each tree level, using the prev. columns
colsample_bynode		Subsamples columns for each split, using the prev. columns


Sample Code




 Dataset   5298

Credit Card Fraud Detection

Anonymized credit card transactions labeled as fraudulent or genuine

 Machine Learning Group - ULB • updated 2 years ago (Version 3)

[Data](#) [Tasks \(9\)](#) [Kernels \(2,449\)](#) [Download \(144 MB\)](#) [New Notebook](#) 

Sample Code

Columns:

- ▶ Time – time from first transaction in dataset
- ▶ V1-V28 – de-identified numeric data
- ▶ Amount – amount in transaction
- ▶ Class – 1 fraudulent, 0 otherwise

Imbalanced classes

Class	
0	99.83%
1	0.17%

Preliminaries

- ▷ Separate data into y (Class) and X (other cols)
- ▷ Split data into train and test sets

Separate the data into X and y

```
In [32]: X = data.iloc[:, :-1]
        y = data.iloc[:, -1]
```


Split the data into train and test sets

```
In [35]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```


Sample Code

`n_estimators = 20, random state = 2020`

Default



max_delta_step = 1





2

XGB Tuning Options

Tuning process & available methods

Methods for Tuning

2 common methods:





Methods for Tuning

Becoming popular: Bayesian Optimization



There are others!



Methods for Tuning

Which one to use?

Method	Speed	No. params	Results
Grid Search	Poor	Poor	Great
Random Search	Great	Great	Good
Bayesian	Good	Great	Great



Tuning Packages in Python

Python

Grid Search

`sklearn.model_selection.GridSearchCV`

Random Search

`sklearn.model_selection.RandomizedSearchCV`

Bayesian

`bayes_opt.BayesianOptimization`

Sample Code: Random Search

Preliminaries

- ▷ Import packages
- ▷ Define parameters
- ▷ Create model (metric = Precision-Recall AUC)

```
In [96]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV  
        from scipy import stats
```

```
In [117]: param = {  
          'objective': 'binary:logistic',  
          'random_state': 2020  
        }
```

```
In [118]: xgb_cv_model = XGBClassifier(**param, nthread = -1, eval_metric = 'aucpr')
```

Sample Code: Random Search

Define distributions to sample from

```
In [119]: cv_params = {  
    'eta': stats.uniform(0.01, 0.59),  
    'min_child_weight': stats.randint(1, 10),  
    'gamma': stats.uniform(0, 10),  
    'subsample': stats.uniform(0.5, 0.5),  
    'colsample_bytree': stats.uniform(0.5, 0.5), → [0.5, 1]  
    'colsample_bylevel': stats.uniform(0.5, 0.5),  
    'colsample_bynode': stats.uniform(0.5, 0.5),  
    'max_depth': [3, 5, 7, 9],  
    'max_delta_step': stats.randint(1, 10)  
}
```

Sample Code: Random Search

Run the random search

```
In [120]: cv_model = RandomizedSearchCV(estimator = xgb_cv_model,  
                                         param_distributions = cv_params,  
                                         random_state = 2020,  
                                         verbose = 2,  
                                         n_jobs = 1,  
                                         )  
cv_model.fit(X_train, y_train)
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits


[Parallel(n_jobs=1)]: Done 30 out of 30 | elapsed: 3.9min finished

```
In [121]: print(cv_model.best_params_)  
  
{'colsample_bylevel': 0.6496611229613486, 'colsample_bynode': 0.834270438868537  
8, 'colsample_bytree': 0.9762762178047335, 'eta': 0.49156213718080255, 'gamma':  
0.7786223582619556, 'max_delta_step': 3, 'max_depth': 5, 'min_child_weight': 7,  
'subsample': 0.9972293498411935}
```

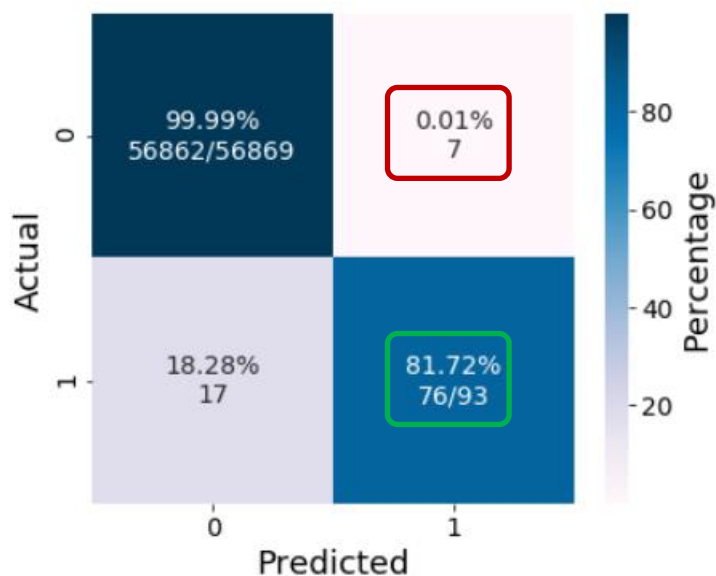
Sample Code: Random Search

`n_estimators = 100, random state = 2020`

No Tuning



With Tuning



Sample Code: Bayesian Optimization

Preliminaries

- ▷ Import packages
- ▷ Create D-matrices (default API easier than sklearn API)

```
In [141]: from bayes_opt import BayesianOptimization
```

```
In [148]: D_train = DMatrix(X_train, label = y_train)
```

```
In [149]: D_test = DMatrix(X_test, label = y_test)
```

Sample Code: Bayesian Optimization

Define upper and lower bounds to sample from

▷ pbounds = dictionary of parameter bounds

```
In [404]: pbounds = {  
    'eta': (0.01, 0.6),  
    'min_child_weight': (1, 10),  
    'gamma': (0, 10),  
    'subsample': (0.5, 1),  
    'colsample_bytree': (0.5, 1),  
    'colsample_bylevel': (0.5, 1),  
    'colsample_bynode': (0.5, 1),  
    'max_depth': (3, 10),  
    'max_delta_step': (1, 10)  
}
```

Sample Code: Bayesian Optimization

■ How to run BayesianOptimization

- ▷ f = function that takes in the parameters being optimized, and outputs a real number (metric: PR AUC of val. set)
- ▷ pbounds – earlier slide

```
xgbc = xgb.cv(  
    paramt,  
    D_train,  
    num_boost_round = 100,  
    nfold = folds,  
    metrics = 'aucpr',  
    stratified = True,  
)
```

Sample Code: Bayesian Optimization

■ Maximize the AUC PR! (took ~1 hour)

```
In [171]: XGB_BO = BayesianOptimization(  
    f=XGB_CV,  
    pbounds=pbounds,  
    random_state=2020,  
    verbose = 10  
)
```

```
In [172]: XGB_BO.maximize(init_points = 15, n_iter = 30)
```


```
In [429]: XGB_BO.max['params']
```

```
Out[429]: {'colsample_bylevel': 0.8429091762081982,  
    'colsample_bynode': 0.5186575195775127,  
    'colsample_bytree': 0.816311624054564,  
    'eta': 0.17053487229722672,  
    'gamma': 1.7892788652395508,  
    'max_delta_step': 9.947823071229202,  
    'max_depth': 9.794647541185961,  
    'min_child_weight': 1.3421690213440391,  
    'subsample': 0.8351187167964789}
```



Sample Code: Bayesian Optimization

`n_estimators = 100, random state = 2020`

RandomSearch



BayesianOptimization



Sample Code: Bayesian Optimization

Comparison of parameters

Parameter	RandomSearch	BayesianOptimization
eta	0.49	0.17
gamma	0.78	1.79
max_delta_step	3	9
max_depth	5	9
min_child_weight	7	1.34
subsample	1.00	0.84
colsample_bylevel	0.65	0.84
colsample_bynode	0.83	0.52
colsample_bytree	0.98	0.82

3

Intro to CatBoost

Categorical Boosting



History of CatBoost

March, 2014

XGBoost initially started
as research project by
Tianqi Chen
but it actually became
famous in 2016

Jan, 2017

Microsoft released
first stable version
of LightGBM

April, 2017

Yandex, one of Russia's
leading tech companies
open sources CatBoost



Comparison of Boosting Algorithms

■ XGBoost

- ▷ Biggest community
- ▷ Most resources

■ LightGBM

- ▷ Fastest
- ▷ The new “meta” for Kaggle competitions
- ▷ You have to convert categorical data to integers



■ CatBoost

- ▷ Newest
- ▷ Smallest community (but has Telegram w/ developers)
- ▷ Best with categorical data, can be used w/o conversion

Features of CatBoost

■ Great performance right out of the box

- ▶ “CatBoost with default parameters beats all other algorithms with tuned parameters except for one case” - Anna Dorogush (Yandex)


	CatBoost		LightGBM		XGBoost		H2O	
	Tuned	Default	Tuned	Default	Tuned	Default	Tuned	Default
 Adult	0.26974	0.27298 +1.21%	0.27602 +2.33%	0.28716 +6.46%	0.27542 +2.11%	0.28009 +3.84%	0.27510 +1.99%	0.27607 +2.35%
 Amazon	0.13772	0.13811 +0.29%	0.16360 +18.80%	0.16716 +21.38%	0.16327 +18.56%	0.16536 +20.07%	0.16264 +18.10%	0.16950 +23.08%




Features of CatBoost

- Tree building is symmetric
 - Each level has the same feature & number
 - You can build trees faster

XGBoost



CatBoost





Features of CatBoost

- Feature interpretations:
 - ▷ Feature importance
 - ▷ Feature interactions
 - ▷ Per object feature importance (SHAP)
- CatBoost Viewer:
 - ▷ Can plot loss function during fitting

Sample Code

Preliminaries

- ▷ Import packages
- ▷ Split data into y (labels) and X (other columns)
- ▷ Train-test split

```
In [212]: from catboost import CatBoostClassifier, Pool, cv
```

Sample Code

Model and plot the evaluation metric

```
In [233]: model = CatBoostClassifier(random_state=2020)
```

```
In [234]: model.fit(X_train, y_train, plot = True)
preds_class = model.predict(X_test)
preds_proba = model.predict_proba(X_test)
```

☒ Learn ☒ Eval

☒ catboost_info

50s 826ms


---- learn

curr ---- 0.0004812...

495

best

Logloss



☐ Click Mode


☒ Logarithm

☐ Smooth


Sample Code

Comparison of results

XGBoost Default



CatBoost Default



Cross Validation with CatBoost

```
In [368]: cv_dataset = Pool(data = X_train,  
                             label = y_train)
```


```
params = {'loss_function': 'RMSE'}
```

```
scores = cv(cv_dataset,  
            params,  
            fold_count = 5,  
            plot = 'True',  
            seed = 2020)
```

```
bestTest = 0.002279584376
```

```
bestIteration = 285
```

```
Shrink model to first 286 iterations.
```





Thank you!

♥ References

- [BayesianOptimization github](#)
- [Bayesian Optimization of XGBoost Parameters](#)
- [CatBoost – Anna Veronika Dorogush \(Yandex\) @ PyData](#)
- [CatBoost GPU Performance](#)
- [CatBoost Docs](#)
- [Details on CatBoost Feature Importances](#)
- [CatBoost: unbiased boosting with categorical features \(paper\)](#)

Special thanks to all the people who made and released these awesome resources for free:

- Presentation template by [SlidesCarnival](#)
- Photographs by [Startup Stock Photos](#)